

Badzioch_Final

December 12, 2024

1 Andrew Badzioch

1.1 Final: Profit Margins

1.1.1 Introduction

- Case Scenario: Smart Electronics Store – Maximizing Profit Margins
- Background: Smart Electronics Store (SES) is a retailer that has rapidly expanded to become a significant player in the electronics market, selling a variety of gadgets and accessories both online and in brick-and-mortar stores. Their product range includes smartphones, charging cables, headphones, monitors, and more niche items like smartwatches and gaming accessories.
- Challenge: SES operates in a highly competitive market where customer preferences and technology trends shift rapidly. The leadership team wants to ensure the company's growth remains robust by maximizing profit margins across its diverse product lines. They seek to understand which products yield the highest profit margins and how they can predict these margins to strategize effectively.
- Data at Hand: SES has collected detailed sales data that includes the cost price, selling price, and quantities sold for their inventory. This dataset is a valuable asset that can be leveraged to predict future profit margins.
- Importance of Margin Prediction:
 1. **Strategic Stocking:** Predicting profit margins will allow SES to prioritize products that offer the best return on investment in terms of shelf space and inventory holdings.
 2. **Dynamic Pricing:** By understanding projected margins, SES can apply dynamic pricing strategies, adjusting prices in real time based on demand, competition, and expected profitability.
 3. **Tailored Promotions:** Products with substantial predicted margins are bundled or promoted to enhance sales volume while maintaining profitability.
 4. **Supplier Negotiations:** With clear insights into which products are likely to be more profitable, SES can negotiate better terms with suppliers or seek cost-effective alternatives.
 5. **Cost Management:** Margin prediction can help identify less profitable products, prompting a review of associated costs or a strategic decision to phase out specific items.
 6. **Budget Allocation:** Accurate margin forecasts are crucial for financial planning, determining where to allocate marketing dollars, and making decisions on product development.

- **Proposed Solution:** To address the challenge, SES plans to use its sales data to build predictive models. By applying algorithms that can handle the variety and complexity of the data, such as Lasso Regression and Support Vector Machines (SVM), SES can predict future profit margins with greater accuracy.
- **Lasso Regression vs. SVM:**
 - **Lasso Regression:** This method is suitable for scenarios where we need to reduce the complexity of a model by performing feature selection. Lasso can identify which product features are most influential on profit margins by penalizing less significant variables to zero.
 - **Support Vector Machines (SVM):** SVM is effective when the relationship between product features and profit margins is not linear or straightforward. Using kernel functions like the radial basis function (RBF), SVM can model complex, non-linear relationships in the data.
- **Conclusion:** By effectively predicting profit margins, SES can make data-driven decisions to streamline their operations, focus on the most profitable products, and adapt their business strategy to the ever-changing market dynamics. This forward-looking approach is essential for maintaining a competitive edge in the technology retail space. Profit margin, also known as net margin, is a financial metric used to assess a company's profitability. It is the percentage of revenue that exceeds the cost of goods sold (COGS), and it is a clear indicator of a company's financial health and efficiency. Essentially, profit margin measures how much of every dollar of sales a company actually keeps in earnings.
- A higher profit margin indicates a more profitable company that has better control over its costs compared to its competitors. Profit margins can vary by industry, and a 'good' margin will often depend on the norms within the particular sector of operation.
- The profit margin for each product can be calculated using the 'Price Each' and 'Cost price' columns from the sales_data.csv file. The 'Price Each' column indicates the selling price of the product, while the 'Cost price' column shows the cost of the product to the business. The difference between these two gives you the gross profit for each product.
- **Data Discription:**
 - **Order Date:** The date and time when the order was placed. Data Type: String (should ideally be datetime).
 - **Order ID:** A unique identifier for the order. Data Type: Integer.
 - **Product:** The name of the product ordered. Data Type: String.
 - **Product_ean:** The European Article Number (EAN), which is a barcode standard, a 12- or 13-digit product identification code. Data Type: String (should be numeric or string due to leading zeroes).
 - **catégorie:** The French word for 'category,' this column likely represents the category of the product. Data Type: String.
 - **Purchase Address:** The address where the product was delivered. Data Type: String.
 - **Quantity Ordered:** The number of units of the product ordered. Data Type: Integer.
 - **Price Each:** The price of one unit of the product. Data Type: Float.
 - **Cost price:** The cost to the company of one unit of the product. Data Type: Float.
 - **Turnover:** This is likely the total revenue from this order line (Quantity Ordered * Price Each). Data Type: Float.

- Margin: The profit margin per product could be the absolute margin (Price Each - Cost price) or a percentage. Data Type: Float.

1.1.2 Importing the libraries needed

```
[1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

import os
```

1.1.3 Loading the Dataset

```
[2]: df = pd.read_csv('sales_data.csv')
df.head()
```

```
[2]:
```

	Order Date	Order ID	Product	Product_ean \
0	1/22/2019 21:25	141234	iPhone	5.638010e+12
1	1/28/2019 14:15	141235	Lightning Charging Cable	5.563320e+12
2	1/17/2019 13:33	141236	Wired Headphones	2.113970e+12
3	1/5/2019 20:33	141237	27in FHD Monitor	3.069160e+12
4	1/25/2019 11:59	141238	Wired Headphones	9.692680e+12

	catégorie	Purchase Address	Quantity Ordered \
0	Vêtements	944 Walnut St, Boston, MA 02215	1
1	Alimentation	185 Maple St, Portland, OR 97035	1
2	Vêtements	538 Adams St, San Francisco, CA 94016	2
3	Sports	738 10th St, Los Angeles, CA 90001	1
4	Électronique	387 10th St, Austin, TX 73301	1

	Price Each	Cost price	turnover	margin
0	700.00	231.0000	700.00	469.0000
1	14.95	7.4750	14.95	7.4750
2	11.99	5.9950	23.98	11.9900
3	149.99	97.4935	149.99	52.4965
4	11.99	5.9950	11.99	5.9950

```
[3]: sales = df.copy()
```

1.1.4 Basic Data Exploration

```
[4]: print("sales dimensions: ", sales.ndim, '\n')
      print("sales shape: ", sales.shape, '\n')
      print("sales size: ", sales.size, '\n')
      print("sales dtype: \n ", sales.dtypes,)
      print("sales columns: ", sales.columns)
```

sales dimensions: 2

sales shape: (185950, 11)

sales size: 2045450

sales dtype:

Order Date	object
Order ID	int64
Product	object
Product_ean	float64
catégorie	object
Purchase Address	object
Quantity Ordered	int64
Price Each	float64
Cost price	float64
turnover	float64
margin	float64

dtype: object

sales columns: Index(['Order Date', 'Order ID', 'Product', 'Product_ean', 'catégorie', 'Purchase Address', 'Quantity Ordered', 'Price Each', 'Cost price', 'turnover', 'margin'], dtype='object')

```
[5]: sales.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 185950 entries, 0 to 185949

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	Order Date	185950 non-null	object
1	Order ID	185950 non-null	int64
2	Product	185950 non-null	object
3	Product_ean	185950 non-null	float64
4	catégorie	185950 non-null	object
5	Purchase Address	185950 non-null	object
6	Quantity Ordered	185950 non-null	int64
7	Price Each	185950 non-null	float64
8	Cost price	185950 non-null	float64

```
9    turnover          185950 non-null float64
10   margin            185950 non-null float64
dtypes: float64(5), int64(2), object(4)
memory usage: 15.6+ MB
```

```
[6]: # lowercase all the column names and replace spaces with underscores
sales.columns = sales.columns.str.lower().str.replace(' ', '_')

# select only columns with string values
string_column = list(sales.dtypes[sales.dtypes == 'object'].index)

# lowercase and replace spaces with underscores in all columns in all string
↳ columns
for col in string_column:
    sales[col] = sales[col].str.lower().str.replace(' ', '_')
```

```
[7]: # checking the column names to see what might be able to be changed, merged, or
↳ dropped
sales.columns
```

```
[7]: Index(['order_date', 'order_id', 'product', 'product_ean', 'catégorie',
        'purchase_address', 'quantity_ordered', 'price_each', 'cost_price',
        'turnover', 'margin'],
        dtype='object')
```

```
[8]: # looking at the format of the order date column
sales['order_date'].head(1)
```

```
[8]: 0    1/22/2019_21:25
      Name: order_date, dtype: object
```

```
[9]: # setting the datetime format to be able to change the columns (separation) with
↳ pd.to_datetime() method
sales['order_dt'] = pd.to_datetime(sales['order_date'], format='%m/%d/%Y_%H:%M')
```

```
[10]: # creating new columns using pandas pd.to_datetime() method
sales['order_year'] = pd.to_datetime(sales['order_dt']).dt.year
sales['order_month'] = pd.to_datetime(sales['order_dt']).dt.month
sales.columns
```

```
[10]: Index(['order_date', 'order_id', 'product', 'product_ean', 'catégorie',
        'purchase_address', 'quantity_ordered', 'price_each', 'cost_price',
        'turnover', 'margin', 'order_dt', 'order_year', 'order_month'],
        dtype='object')
```

```
[11]: # dropping unneeded columns
sales.drop(columns={'order_date', 'order_id', 'product_ean',
↳ 'purchase_address', 'order_dt'}, axis=1, inplace=True)
```

```
# changing 'catégorie' to 'category'
sales = sales.rename(columns={'catégorie':'category'})

sales.columns
```

```
[11]: Index(['product', 'category', 'quantity_ordered', 'price_each', 'cost_price',
          'turnover', 'margin', 'order_year', 'order_month'],
          dtype='object')
```

```
[12]: sales.head(10)
```

```
[12]:
```

	product	category	quantity_ordered	price_each	\
0	iphone	vêtements	1	700.00	
1	lightning_charging_cable	alimentation	1	14.95	
2	wired_headphones	vêtements	2	11.99	
3	27in_fhd_monitor	sports	1	149.99	
4	wired_headphones	électronique	1	11.99	
5	aaa_batteries_(4-pack)	alimentation	1	2.99	
6	27in_4k_gaming_monitor	vêtements	1	389.99	
7	usb-c_charging_cable	vêtements	1	11.95	
8	bose_soundsport_headphones	électronique	1	99.99	
9	apple_airpods_headphones	électronique	1	150.00	

	cost_price	turnover	margin	order_year	order_month
0	231.0000	700.00	469.0000	2019	1
1	7.4750	14.95	7.4750	2019	1
2	5.9950	23.98	11.9900	2019	1
3	97.4935	149.99	52.4965	2019	1
4	5.9950	11.99	5.9950	2019	1
5	1.4950	2.99	1.4950	2019	1
6	128.6967	389.99	261.2933	2019	1
7	5.9750	11.95	5.9750	2019	1
8	49.9950	99.99	49.9950	2019	1
9	97.5000	150.00	52.5000	2019	1

```
[13]: sales['product'].value_counts()
```

```
[13]: product
usb-c_charging_cable      21903
lightning_charging_cable  21658
aaa_batteries_(4-pack)    20641
aa_batteries_(4-pack)     20577
wired_headphones          18882
apple_airpods_headphones  15549
bose_soundsport_headphones 13325
27in_fhd_monitor          7507
```

```

iphone                                6842
27in_4k_gaming_monitor                6230
34in_ultrawide_monitor                6181
google_phone                          5525
flatscreen_tv                         4800
macbook_pro_laptop                    4724
thinkpad_laptop                       4128
20in_monitor                          4101
vareebadd_phone                       2065
lg_washing_machine                    666
lg_dryer                              646
Name: count, dtype: int64

```

```

[14]: # changing the values of category to make more effecient
sales.loc[sales['product'].isin(['iphone', 'google_phone', 'vareebadd_phone']),
        ↪ 'category'] = 'phone'
sales.loc[sales['product'].isin(['aaa_batteries_(4-pack)',
        ↪ 'aa_batteries_(4-pack)']), 'category'] = 'batteries'
sales.loc[sales['product'].isin(['wired_headphones',
        ↪ 'apple_airpods_headphones', 'bose_soundsport_headphones']), 'category'] =
        ↪ 'headphones'
sales.loc[sales['product'].isin(['macbook_pro_laptop', 'thinkpad_laptop']),
        ↪ 'category'] = 'laptop'
sales.loc[sales['product'].isin(['27in_fhd_monitor', '27in_4k_gaming_monitor',
        ↪ '34in_ultrawide_monitor', 'flatscreen_tv', '20in_monitor']), 'category'] =
        ↪ 'monitor'
sales.loc[sales['product'].isin(['usb-c_charging_cable',
        ↪ 'lightning_charging_cable']), 'category'] = 'cables'
sales.loc[sales['product'].isin(['lg_washing_machine', 'lg_dryer']),
        ↪ 'category'] = 'misc'

sales.head(10)

```

```

[14]:
      product  category  quantity_ordered  price_each  \
0         iphone      phone                1         700.00
1  lightning_charging_cable    cables                1          14.95
2         wired_headphones  headphones                2          11.99
3         27in_fhd_monitor    monitor                1         149.99
4         wired_headphones  headphones                1          11.99
5    aaa_batteries_(4-pack)  batteries                1           2.99
6    27in_4k_gaming_monitor    monitor                1        389.99
7         usb-c_charging_cable    cables                1          11.95
8  bose_soundsport_headphones  headphones                1          99.99
9  apple_airpods_headphones  headphones                1        150.00

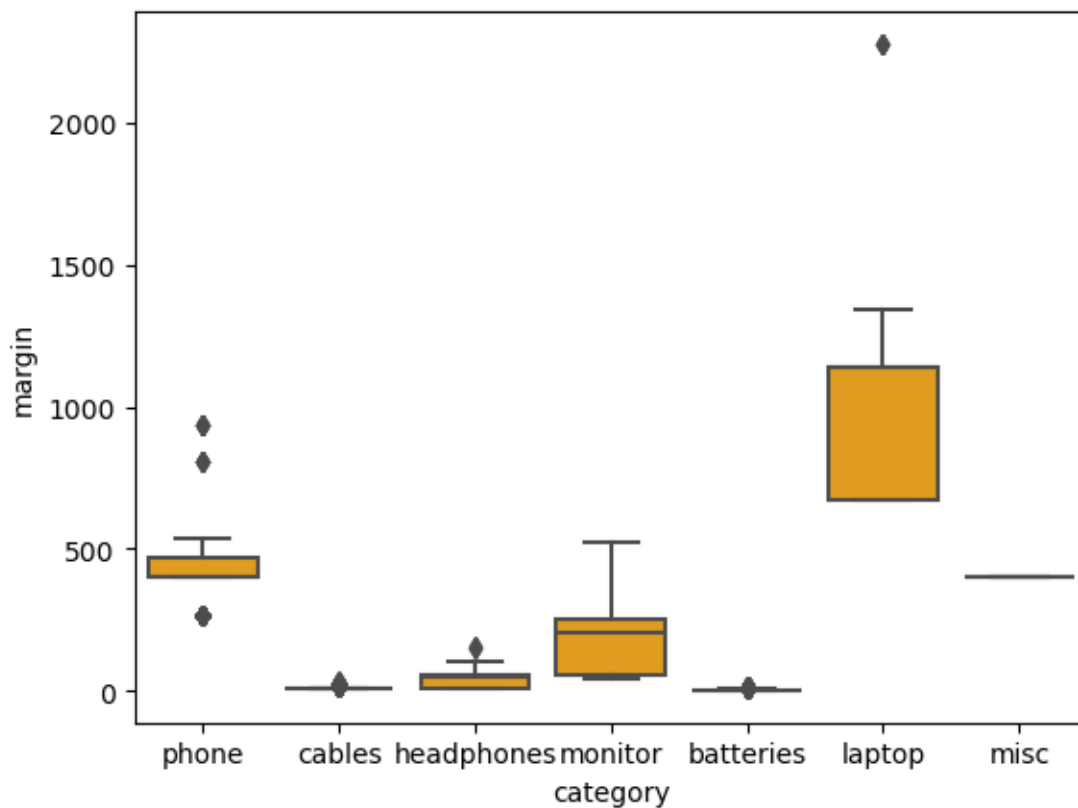
      cost_price  turnover   margin  order_year  order_month
0      231.0000    700.00  469.0000        2019            1

```

1	7.4750	14.95	7.4750	2019	1
2	5.9950	23.98	11.9900	2019	1
3	97.4935	149.99	52.4965	2019	1
4	5.9950	11.99	5.9950	2019	1
5	1.4950	2.99	1.4950	2019	1
6	128.6967	389.99	261.2933	2019	1
7	5.9750	11.95	5.9750	2019	1
8	49.9950	99.99	49.9950	2019	1
9	97.5000	150.00	52.5000	2019	1

```
[15]: sns.boxplot(x='category', y='margin', data=sales, color='orange')
```

```
[15]: <Axes: xlabel='category', ylabel='margin'>
```



1.2 Case Scenario: Maximizing Profit Margin

```
[16]: # creating a new dataframe for this scenario
sales1 = sales.copy()
sales1.head()
```



```
[16]:
```

	product	category	quantity_ordered	price_each	\
0	iphone	phone	1	700.00	
1	lightning_charging_cable	cables	1	14.95	
2	wired_headphones	headphones	2	11.99	
3	27in_fhd_monitor	monitor	1	149.99	
4	wired_headphones	headphones	1	11.99	

	cost_price	turnover	margin	order_year	order_month
0	231.0000	700.00	469.0000	2019	1
1	7.4750	14.95	7.4750	2019	1
2	5.9950	23.98	11.9900	2019	1
3	97.4935	149.99	52.4965	2019	1
4	5.9950	11.99	5.9950	2019	1

```
[17]: # creating a new column to check the total revenue
sales1['pro_mar'] = (sales1['price_each'] - sales1['cost_price'])
total_pro = sales1['pro_mar'].sum()
print("The total sales revenue is: ", f"{total_pro:.2f}")
```

The total sales revenue is: 21334257.74

```
[18]: # checking that the turnover is the same
print(f"the total margin is: {sales1['margin'].sum():.2f}")
```

the total margin is: 21438067.93

```
[19]: # checking if there is a difference between 'pro_mar' and 'margin'
diff = (sales1['margin'] != sales1['pro_mar']).sum()
print(f"There are {diff} 'margin' and 'pro_mar' rows that do not equal each_
↳ other.")
```

There are 35204 'margin' and 'pro_mar' rows that do not equal each other.

```
[20]: print("the margin of error is: ", f"{abs(total_pro - sales1['margin'].sum()):.
↳ 2f}")
```

the margin of error is: 103810.19

1.2.1 Importing the Libraries for Regression Analysis

```
[21]: from sklearn.model_selection import train_test_split, GridSearchCV,
↳ RandomizedSearchCV
from sklearn.linear_model import LassoCV
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import OneHotEncoder, StandardScaler
```

```
[22]: sales1 = sales1.drop(columns='margin')
```

```
[23]: # making sure there are no empty cells (double checking)
sales1.isnull().sum()
```

```
[23]: product          0
      category        0
      quantity_ordered 0
      price_each       0
      cost_price       0
      turnover         0
      order_year       0
      order_month      0
      pro_mar          0
      dtype: int64
```

```
[24]: # defining features and target variable
X = sales1.drop('pro_mar', axis=1)
y = sales1['pro_mar']

# splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[25]: # use pd.get_dummies to one_hot encode categorical features
X_train_encoded = pd.get_dummies(X_train[['product', 'category']],
↳drop_first=True)
X_test_encoded = pd.get_dummies(X_test[['product', 'category']],
↳drop_first=True)
```

```
[26]: X_train_encoded, X_test_encoded = X_train_encoded.align(X_test_encoded,
↳join='left', axis=1)
```

```
[27]: # scaling the numerical columns (only the numerical features)
numb_cols = ['quantity_ordered', 'price_each', 'cost_price', 'order_year',
↳'order_month']
scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train[numb_cols])
X_test_scaled = scaler.transform(X_test[numb_cols])
```

```
[28]: X_train_scaled_df = pd.DataFrame(X_train_scaled, columns=numb_cols,
↳index=X_train.index)
X_test_scaled_df = pd.DataFrame(X_test_scaled, columns=numb_cols, index=X_test.
↳index)

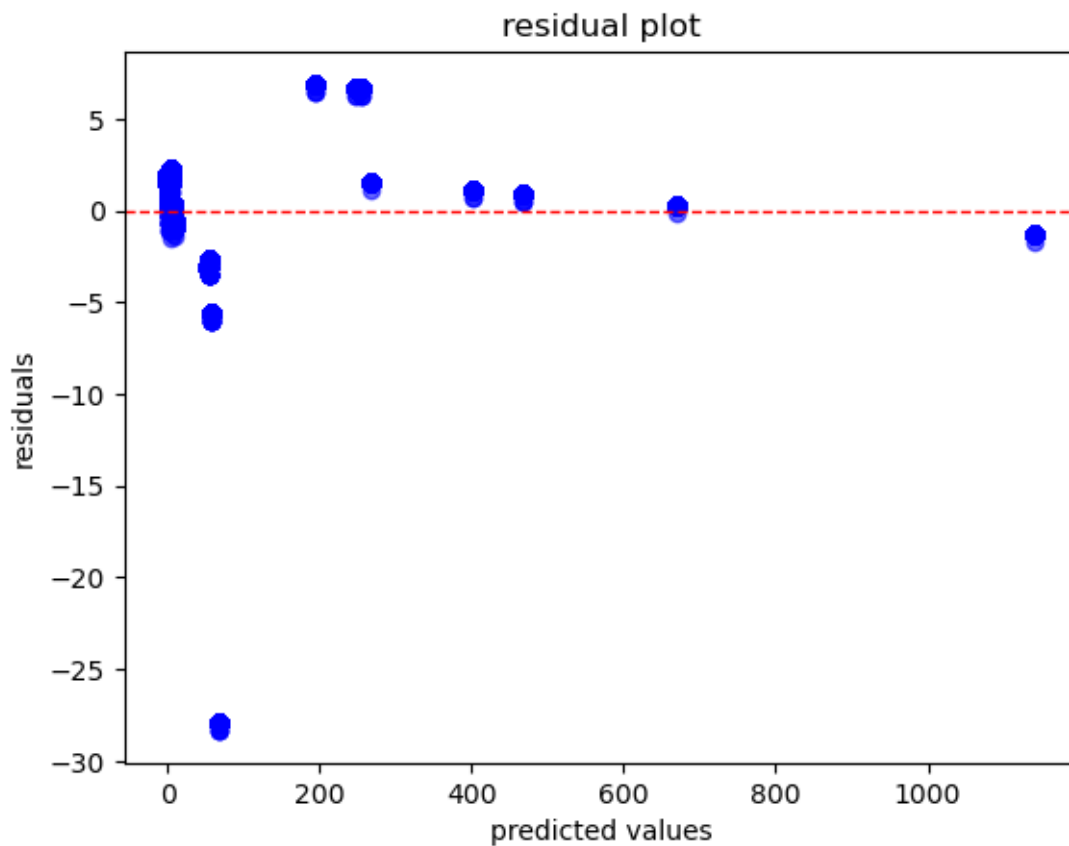
X_train_final = pd.concat([X_train_encoded, X_train_scaled_df], axis=1)
X_test_final = pd.concat([X_test_encoded, X_test_scaled_df], axis=1)
```

```
[29]: lasso = LassoCV(cv=5, random_state=0)
lasso.fit(X_train_final, y_train)
```

```
[29]: LassoCV(cv=5, random_state=0)
```

```
[30]: y_pred_lasso = lasso.predict(X_test_final)
```

```
[31]: # plotting for visualization
residuals = y_test - y_pred_lasso
plt.scatter(y_pred_lasso, residuals, color='blue', alpha=0.6)
plt.axhline(0, color='red', linestyle='--', linewidth=1)
plt.xlabel('predicted values')
plt.ylabel('residuals')
plt.title('residual plot')
plt.show();
```



```
[32]: print(f"Lasso Regression RMSE: {np.sqrt(mean_squared_error(y_test, y_pred_lasso))}")
print(f"Lasso Regression R^2: {r2_score(y_test, y_pred_lasso)}")
```

Lasso Regression RMSE: 4.922092566682225
Lasso Regression R²: 0.999521945406097

```
[33]: # Sample a subset of data
X_train_sample = X_train_scaled[:1000]
y_train_sample = y_train[:1000]
```

```
[34]: # Parameter grid
param_grid = {
    'C': [0.1, 1, 10],
    'gamma': ['scale', 'auto']
}
```

```
[35]: # Randomized search
random_search = RandomizedSearchCV(
    SVR(kernel='rbf'),
    param_distributions=param_grid,
    n_iter=5, # Test fewer combinations
    cv=3,
    n_jobs=-1 # Parallelize
)
random_search.fit(X_train_sample, y_train_sample)
```


```
[35]: RandomizedSearchCV(cv=3, estimator=SVR(), n_iter=5, n_jobs=-1,
                        param_distributions={'C': [0.1, 1, 10],
                        'gamma': ['scale', 'auto']})
```

```
[36]: # Evaluate best model
best_svm = random_search.best_estimator_
y_pred_svm = best_svm.predict(X_test_scaled)
print(f"SVM Regression RMSE: {np.sqrt(mean_squared_error(y_test, y_pred_svm))}")
print(f"SVM Regression R2: {r2_score(y_test, y_pred_svm)}")
```

SVM Regression RMSE: 114.86085466887933
SVM Regression R²: 0.7396713123473979

```
[37]: # plotting output

plt.figure(figsize=(8, 8))
plt.scatter(y_test, y_pred_svm, alpha=0.5)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.xlabel("True Values")
plt.ylabel("Predictions")
plt.title("SVR Predictions vs. True Values")
plt.grid()
plt.show()
```



[]: